



15-110 PRINCIPLES OF COMPUTING – S21

LECTURE 1: INTRODUCTION AND LOGISTICS

TEACHER:
GIANNI A. DI CARO

The 110 team

Teacher:



Prof. Gianni A. Di Caro
(robotics, AI, ML, MAS)

Course Assistants:

- Samar Rahmouni
- Yusuf Saquib
- Jawaria Abbas
- Demetre Devidze
- Abdalla Mohamed

Road map

- General overview: what this course is about
- Utility of the course
- Logistics: classes, labs, homework, exams, grading, website, books, piazza, rules
- Software for python programming
- More about the need to use efficient algorithms + computers

Road map

- General overview: what this course is about
- Utility of the course
- Logistics: classes, labs, homework, exams, grading, website, books, piazza, rules
- Software for python programming
- More about the need to use efficient algorithms + computers

Introduction to the *science (art)* of computational problem solving

Daily life, professional activities, leisure, business, ... present a variety of **problems** that ask for finding a **solution**



- Different **constraints** (e.g., time, budget)
- Different representations / models / **contexts**
- Different **objectives** (e.g., best nutritional balance, minimum traveling time, win!)
- Different **requirements for the solution** (e.g., provably the best, just one)
- Different **available knowledge** (e.g., chess vs. financial investment vs. path finding)
- Different **degree of (un)certainty** (e.g., poker vs. industrial manipulator)
- Different **dimensionality / number of variables** (e.g., DNA sequencing, root finding, flight control)
- ...

Introduction to the *science (art)* of computational problem solving

Often (usually) these problems are too complex to be **effectively** and **satisfactorily** solved by humans (alone)...

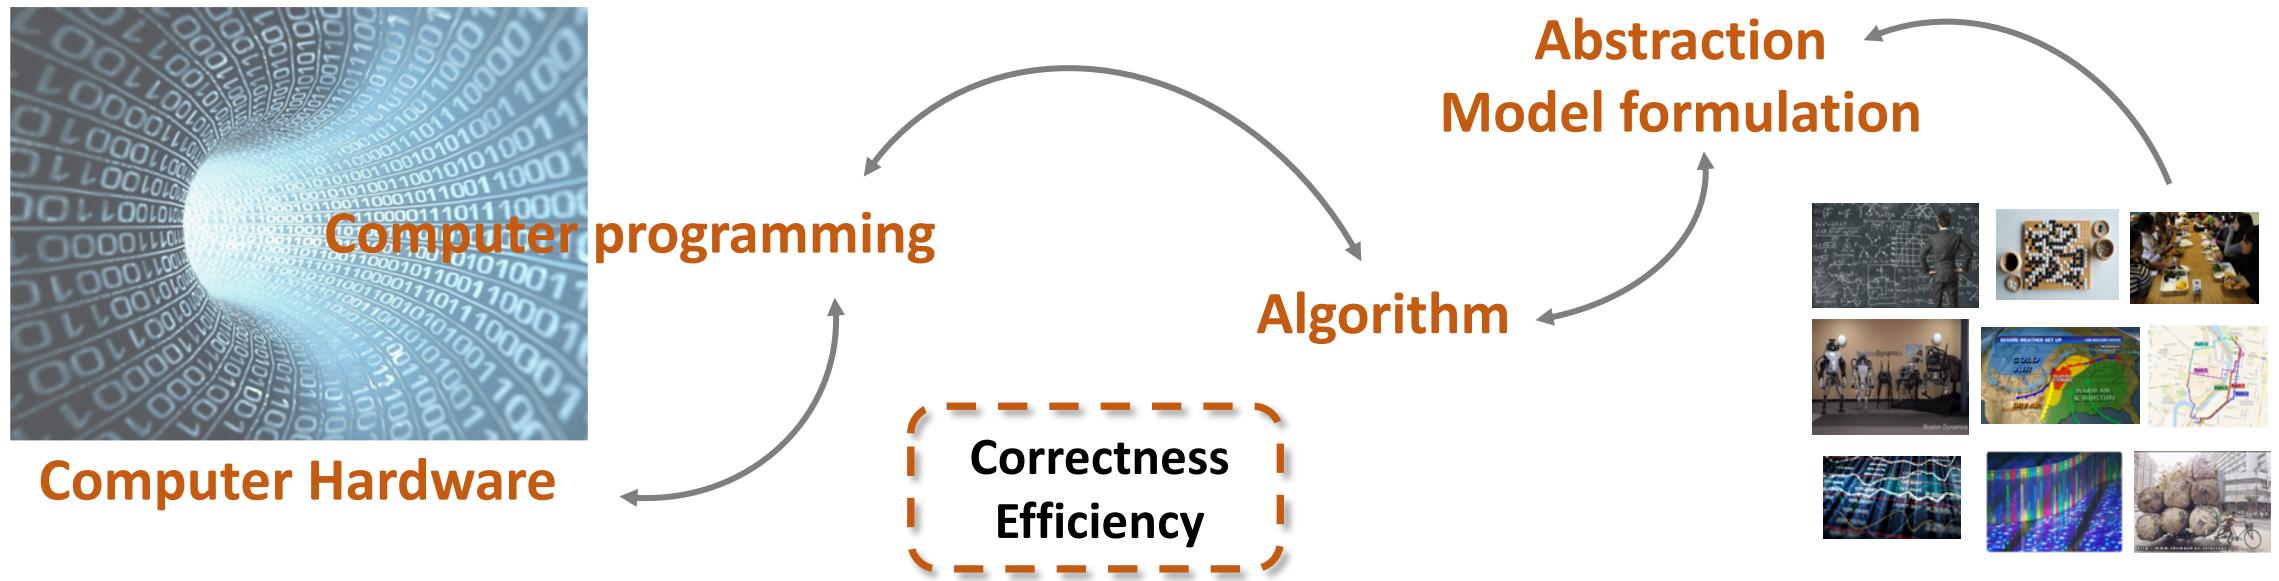
... **computers** can effectively help (us) solving the difficult problems!



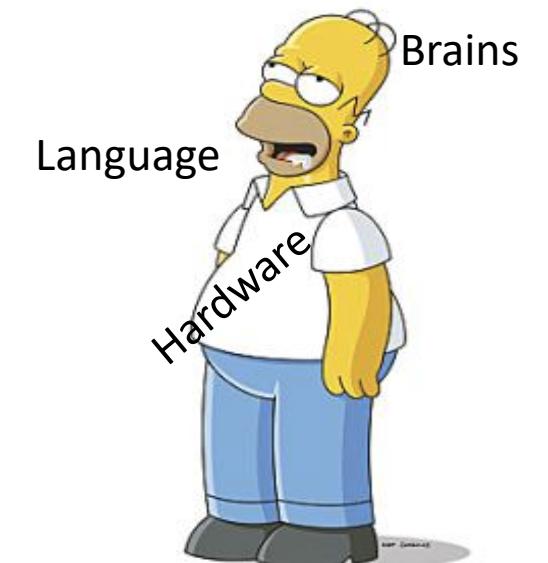
- Perform **calculations**, at impressively high rates
 - + - * / > < = ≠
- Fastest computer in the world (IBM AC922): $\approx 10^{15}$ operations per sec on real numbers:
1,000,000,000,000,000 ops/sec
- Mega = 10^6 , Giga = 10^9 , Tera = 10^{12} , Peta = 10^{15} , Exa = 10^{18}
- **Store** and **retrieve** data and results of calculations
 - 160TB internal memory, HP: 160,000,000,000,000 bytes
 - Single storage units > 100 TB: 100,000,000,000,000 bytes

... how do we connect *problem + humans + computers* for effective problem-solving?

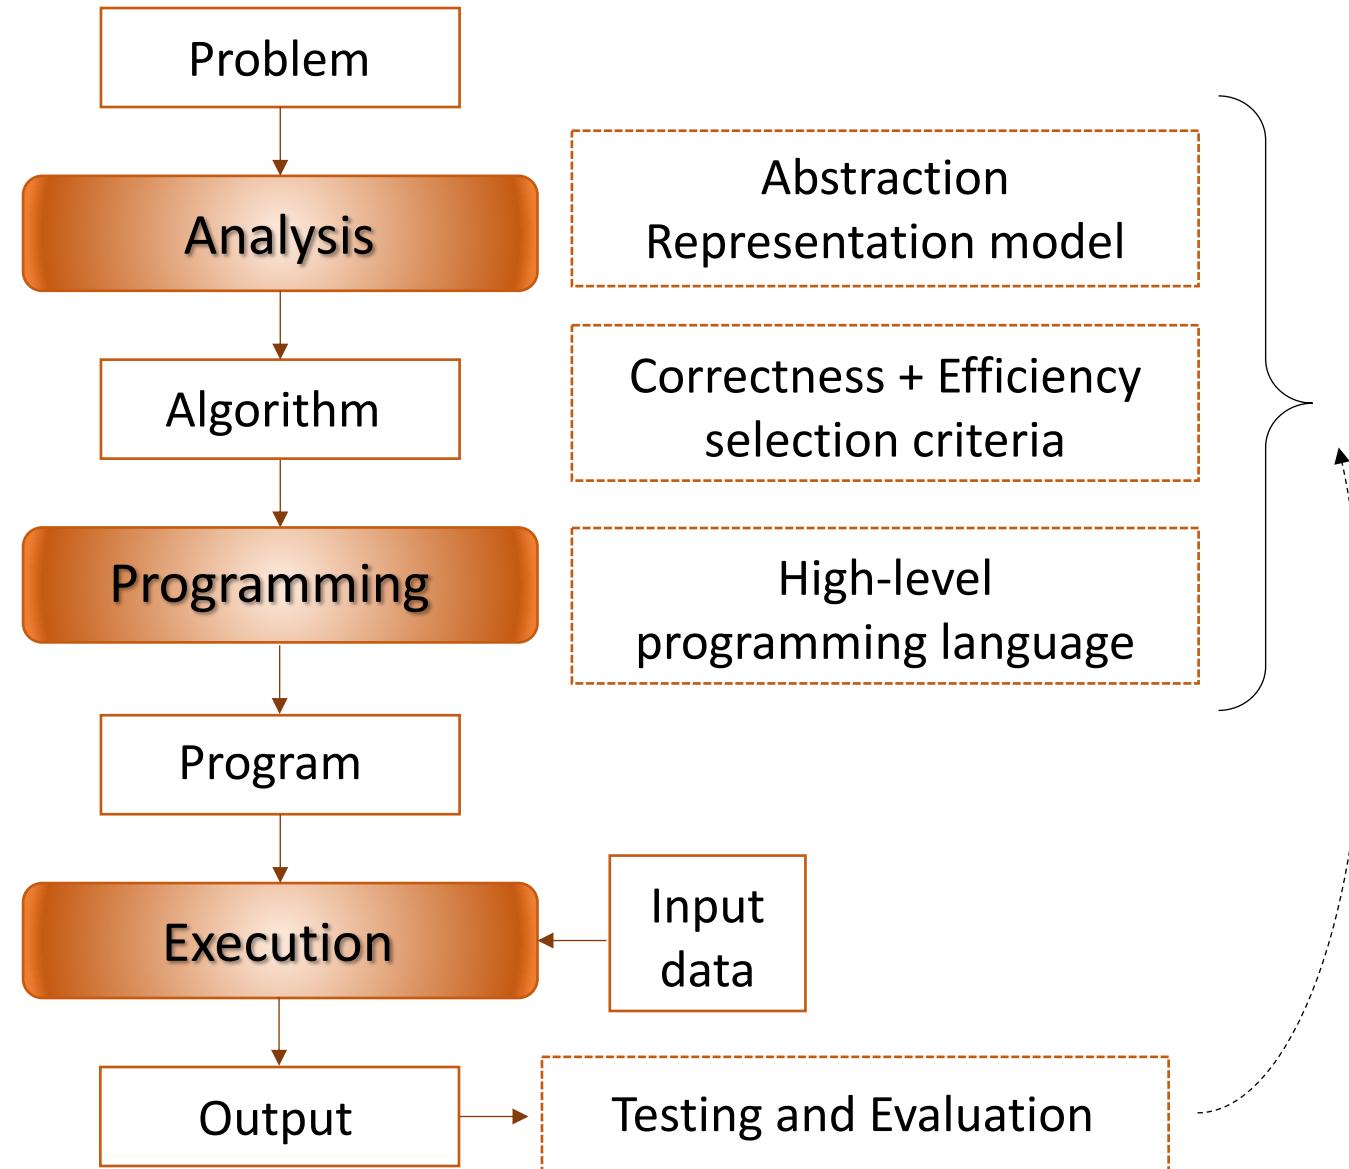
Introduction to the *science (art)* of computational problem solving



- ✓ **Algorithm:** What to do (*step by step*) to (efficiently) solve the problem based on the defined problem abstraction / model → **Problem-solving recipe**
- ✓ **Programming (coding):** Use a high-level (computer) *language* to precisely describe (code) the algorithm → The code is used to tell the computer to do the things prescribed by the algorithm [**what to do and how**]
- ✓ **Computer hardware:** The high-level program is translated into a *machine-level* language which is then **executed** by the specific computer hardware



Introduction to the *science (art)* of computational problem solving



- **Algorithm:**

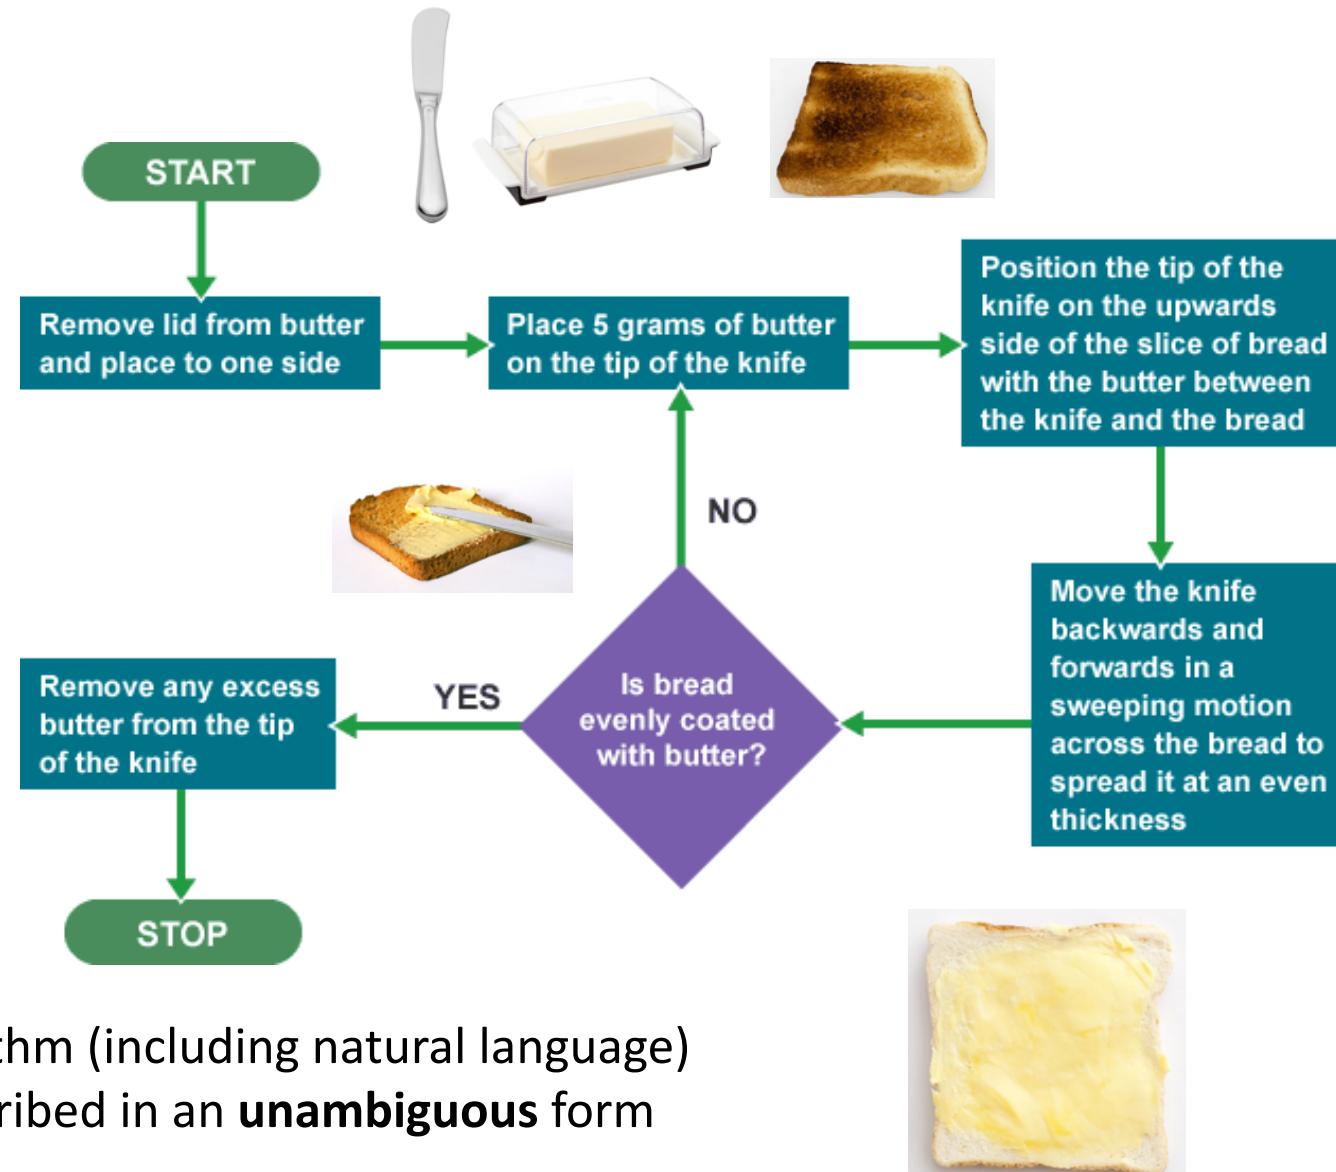
- ✓ A finite list of instructions that describe a **computation**
- ✓ when the instructions are executed on a provided set of inputs, the computation proceeds *step by step* through a set of well-defined states (configurations)
- ✓ eventually, it ends, with some outputs being produced

- **Program:**

- ✓ Algorithm encoding using a language that the computer understands
- ✓ > 700 *programming languages!*
- ✓ Primitive constructs, syntax, static semantics, semantics

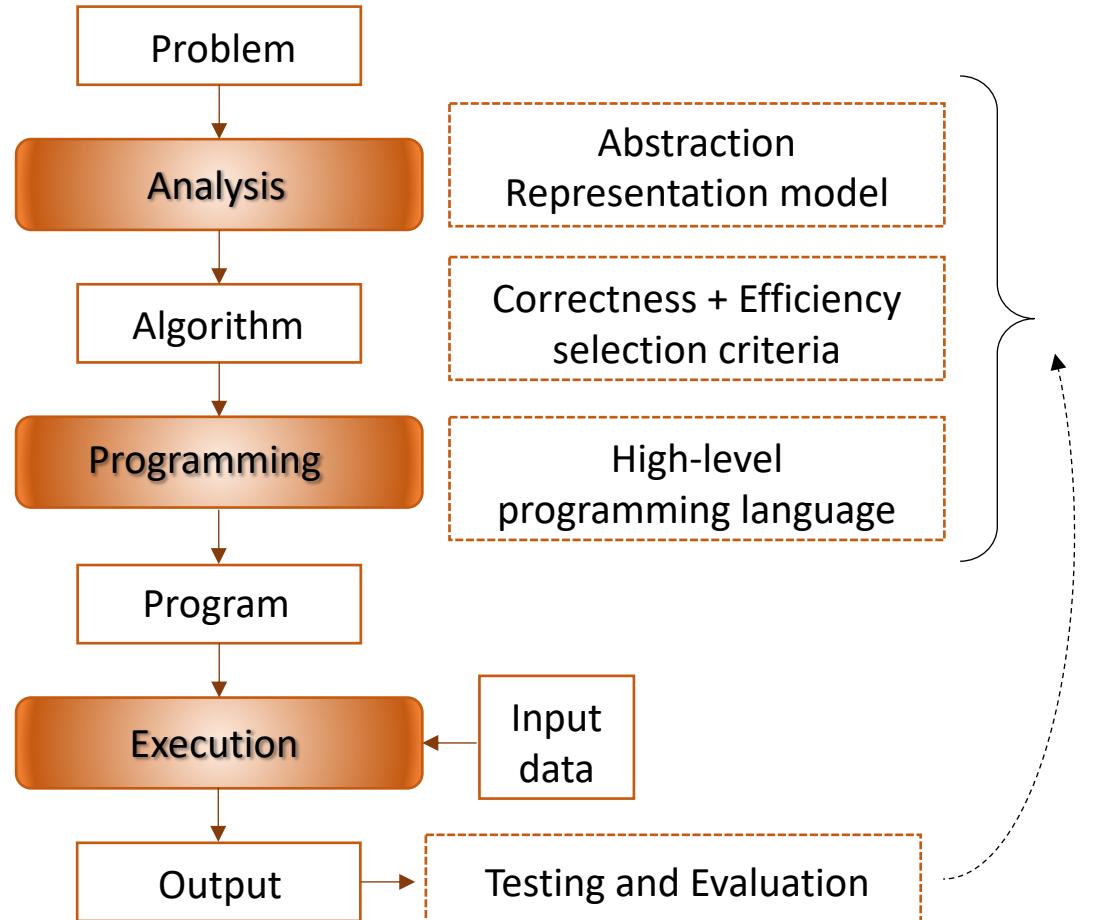
Algorithms

- ✓ A finite list of instructions that describe a **computation**
- ✓ when the instructions are executed on a provided set of inputs, the computation proceeds step by step through a set of well-defined states (configurations)
- ✓ eventually, it ends, with some outputs being produced



- Many ways to **describe** the same algorithm (including natural language)
- To be implemented, it needs to be described in an **unambiguous** form
- Can be applied to a **class of problems**

The 15-110 journey ...



- ✓ **Problem analysis:** abstraction and representation
- ✓ Concepts central to design and understand **computation** (i.e., *algorithms*)
- ✓ Fundamental concepts of **programming**:
 - *Knowledge representation* (data types & structures)
 - *Flow control* (conditionals, iteration, recursion)
 - *Data organization* (lists, matrices, dictionaries)
 - *Decomposition, reusability, information hiding* (modules, functions, objects)
- ✓ **Python programming language**
- ✓ Computational **problem-solving techniques**
- ✓ **Correctness and efficiency** of algorithms (testing, error finding, complexity, optimization)
- ✓ **Data visualization and data publishing**
- ✓ **Put all together into a data science project!**

Road map

- General overview: what this course is about
- **Utility of the course**
- Logistics: classes, labs, homework, exams, grading, website, books, piazza, rules
- Software for python programming
- More about the need to use efficient algorithms + computers

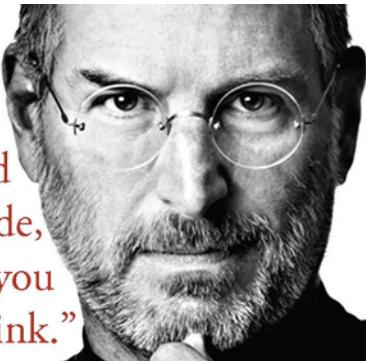
Motivations for following 15-110 ...

- Understanding computation and learning computer programming helps shaping your mind, develops your critical thinking: it's like *math* but more interactive and more fun!

Steve Jobs

1955-2011

“Everyone should
learn how to code,
it teaches you
how to think.”



- ✓ Rational thinking for (general) problem analysis and solving
- ✓ Develop abstraction and formalization capabilities
- ✓ Devising an efficient solution to a difficult problem is always rewarding
- ✓ Finding errors in the code is as thrilling as finding the assassin! ☺

- Why understanding computation and learning computer programming can be rewarding?

- Financial analysts: **\$84,300**
- Market research analysts: **\$63,230**
- Sales managers: **\$121,060**
- Human resource specialists: **\$60,350**
- Accountants and auditors: **\$69,350**
- Management analysts: **\$82,450**
- Personal financial advisors: **\$90,640**

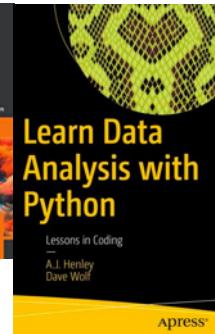
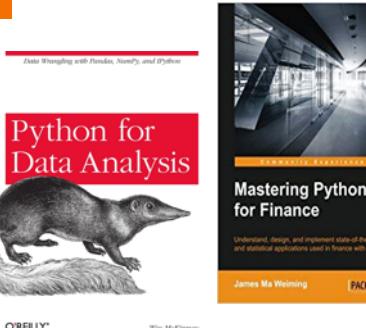
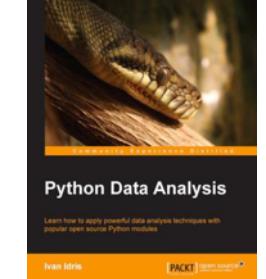
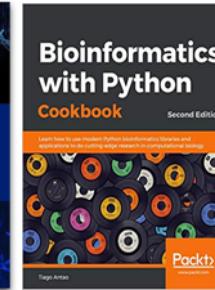
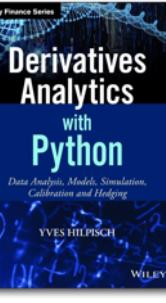
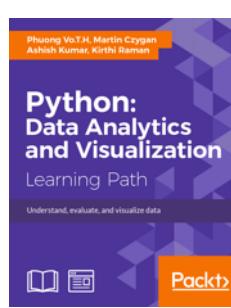
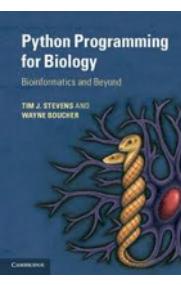
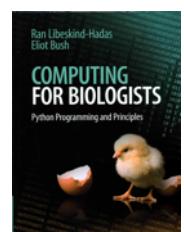
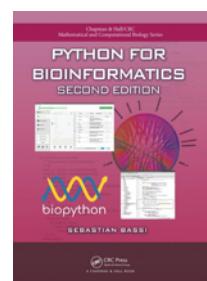
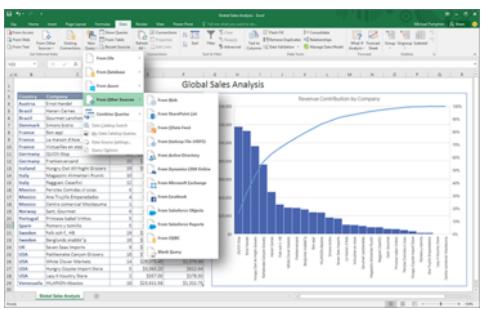
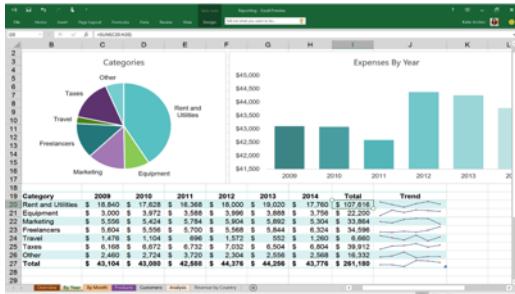
- **Software Engineer:** Average Salary \$76,194
- **Senior Software Engineer:** Average Salary \$89,599
- **Software Developer:** Average Salary \$62,803
- **Senior Software Developer or Programmer:** Average Salary \$100,303
- **Web Developer:** Average Salary \$62,500
- **Programmer Analyst:** Average Salary \$55,250

Salary range for a **Biologist**

\$44,946 to \$56,017

Motivations for following 15-110 ...

- Why understanding computation and learning computer programming can result in a professional advantage also for non-CS professionals?



Road map

- General overview: what this course is about
- Utility of the course
- Logistics: classes, labs, homework, exams, grading, website, books, piazza, rules
- Software for python programming
- More about the need to use efficient algorithms + computers

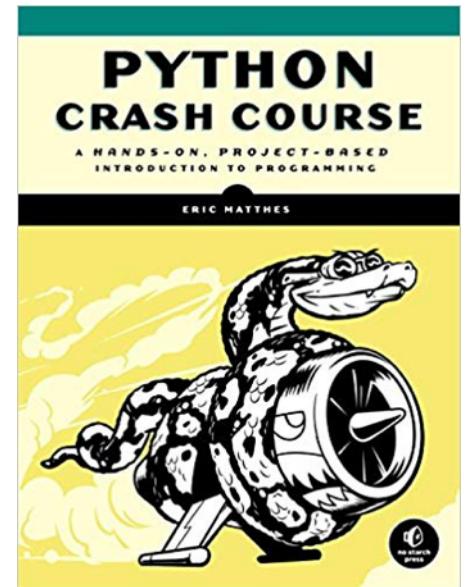
Logistics: Classes and Grading

- **Theory Classes:** Sundays, Tuesdays (1:30pm – 2:50pm), ONLINE
 - **Labs, hands-on, Quizzes:** Thursdays (1:30pm – 2:50pm), room 1064, In-presence
 - **Homework:** weekly, programming questions (autograded) + theory quizzes
 - **Quizzes:** weekly, during lab hours
 - **Midterms:** three tests during the course
 - **Final exam:** one general test at the end of course
 - **Project:** programming project issued at week 10 to be completed by the end of the course (deliver: code/notebook + oral presentation to teacher)
- **Homeworks: 15%**
 - Weekly, except on exam weeks
 - **To be solved individually** (check the course policy)
 - Check the schedule for dates and deadlines
 - **Quizzes: 25%**
 - Weekly
 - Administered at the beginning of the lecture or lab
 - Individual and closed book
 - **Three midterms: 30% (10% each)**
 - Administered during normal class hours
 - Individual and closed book
 - **Project: 10%**
 - **Final exam: 20%**

Logistics: Resources

- Main sources for material:

- Lectures + lecture slides + technical notes
- Python reference book on **Canvas**
- Course website: <https://cs-cmuq.github.io/110-www/>



- **Piazza** for: Questions, announcements, discussions

A screenshot of a web browser displaying the Piazza platform. The address bar shows the URL "piazza.com/class/kk06ethmf7w3r8?cid=7". The main header has the word "piazza" in white on a blue background. Below the header, there are several navigation links: "15-110Q ▾", "Setup", "Q & A" (which is underlined to indicate it's the active tab), "Resources", "Statistics ▾", and "Manage Class". At the bottom of the page, there is a horizontal menu with categories: "LIVE Q&A", "Drafts", "hw1", "hw2", "hw3", "hw4", "hw5", "hw6", "hw7", "hw8", "hw9", "hw10", "project", "exam", "logistics", and "other".

- **Autolab** for submitting homework and project, and getting grades

A screenshot of a web browser displaying the Autolab platform. The address bar shows the URL "autolab.andrew.cmu.edu/courses/15110q-s21/assessments". The main header has the word "AUTOLAB" in white on a red background. Below the header, there is a navigation bar with icons for "Home" and "15110q-s21: Principles of Computing (s21)".

Logistics: In-class rules

Be *alive* and **participate!**



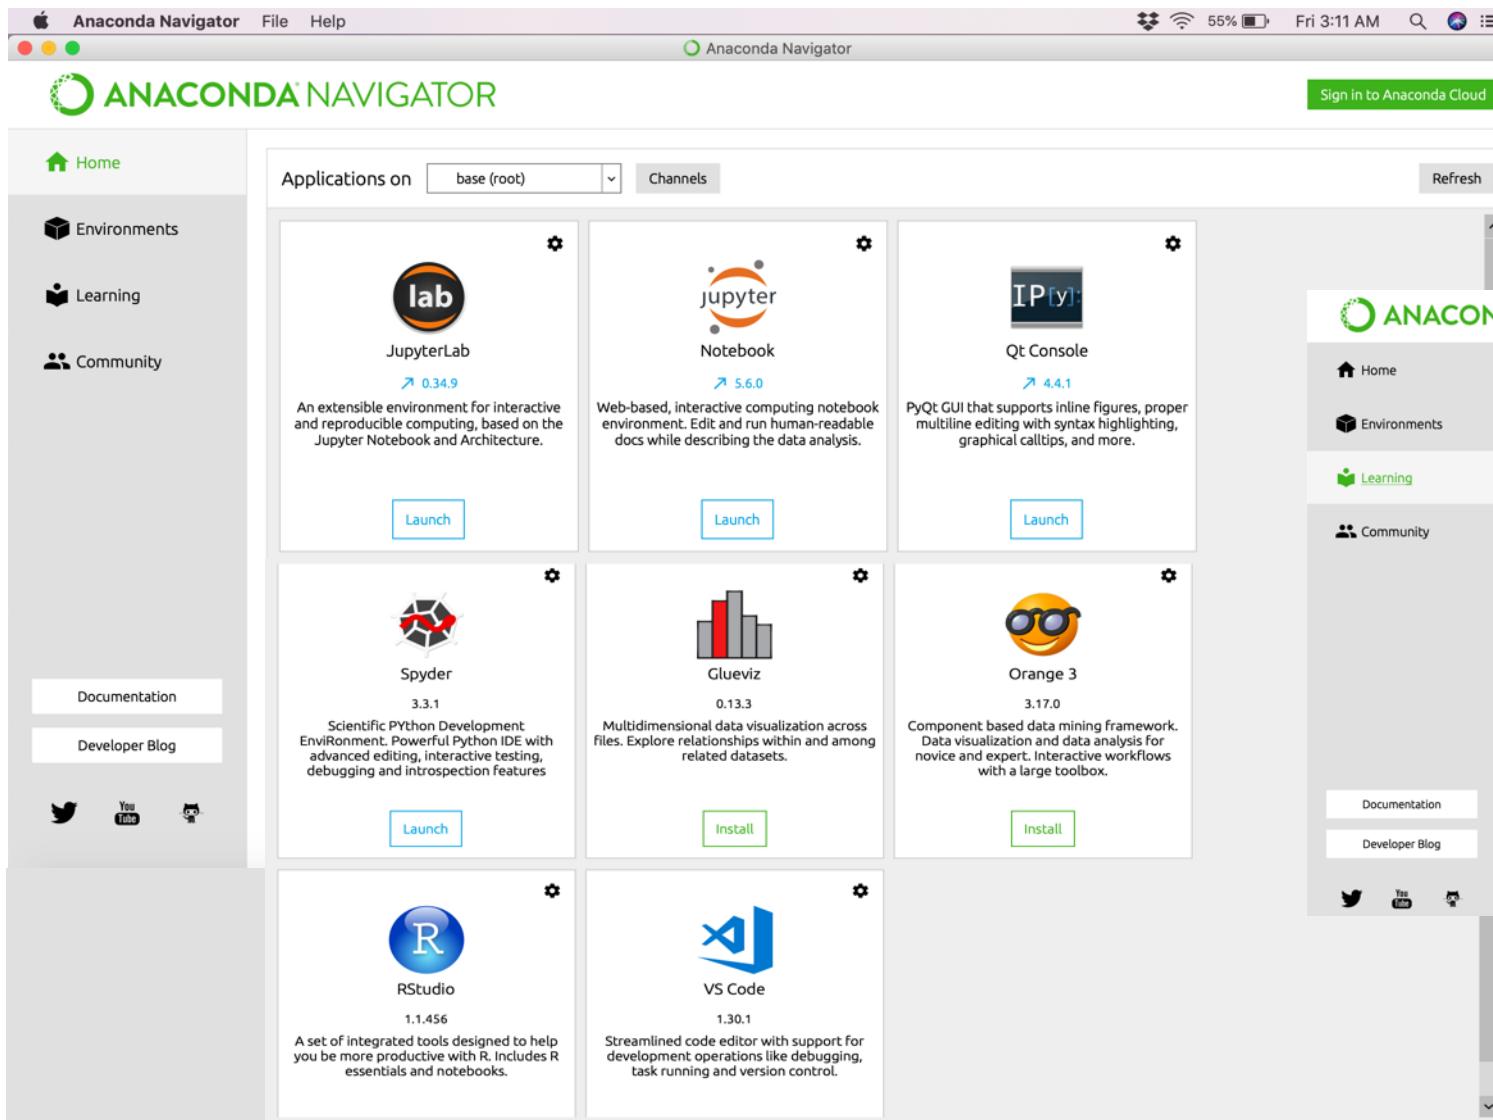
Ask questions!



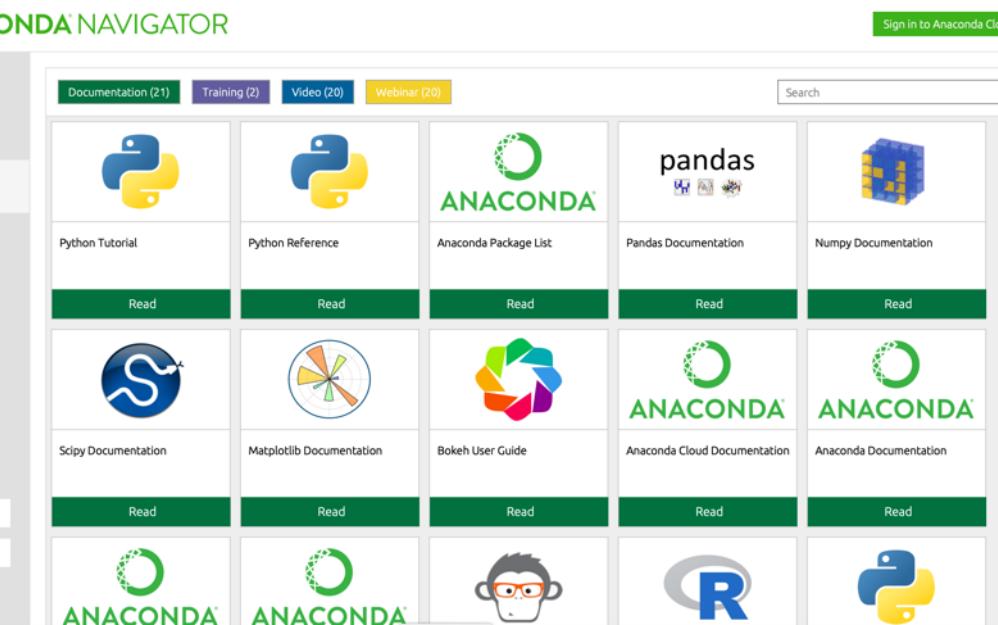
Road map

- General overview: what this course is about
- Utility of the course
- Logistics: classes, labs, homework, exams, grading, website, books, piazza, rules
- **Software for python programming**
- More about the need to use efficient algorithms + computers

ANACONDA Navigator



<https://www.anaconda.com/>



ANACONDA: Install

anaconda.com/products/individual



ANACONDA

Products ▾

Pricing

Solutions ▾

Resources ▾

Blog

Company ▾

Anaconda Installers



Python 3.8

64-Bit Graphical Installer (457 MB)

32-Bit Graphical Installer (403 MB)



Python 3.8

64-Bit Graphical Installer (435 MB)

64-Bit Command Line Installer (428 MB)

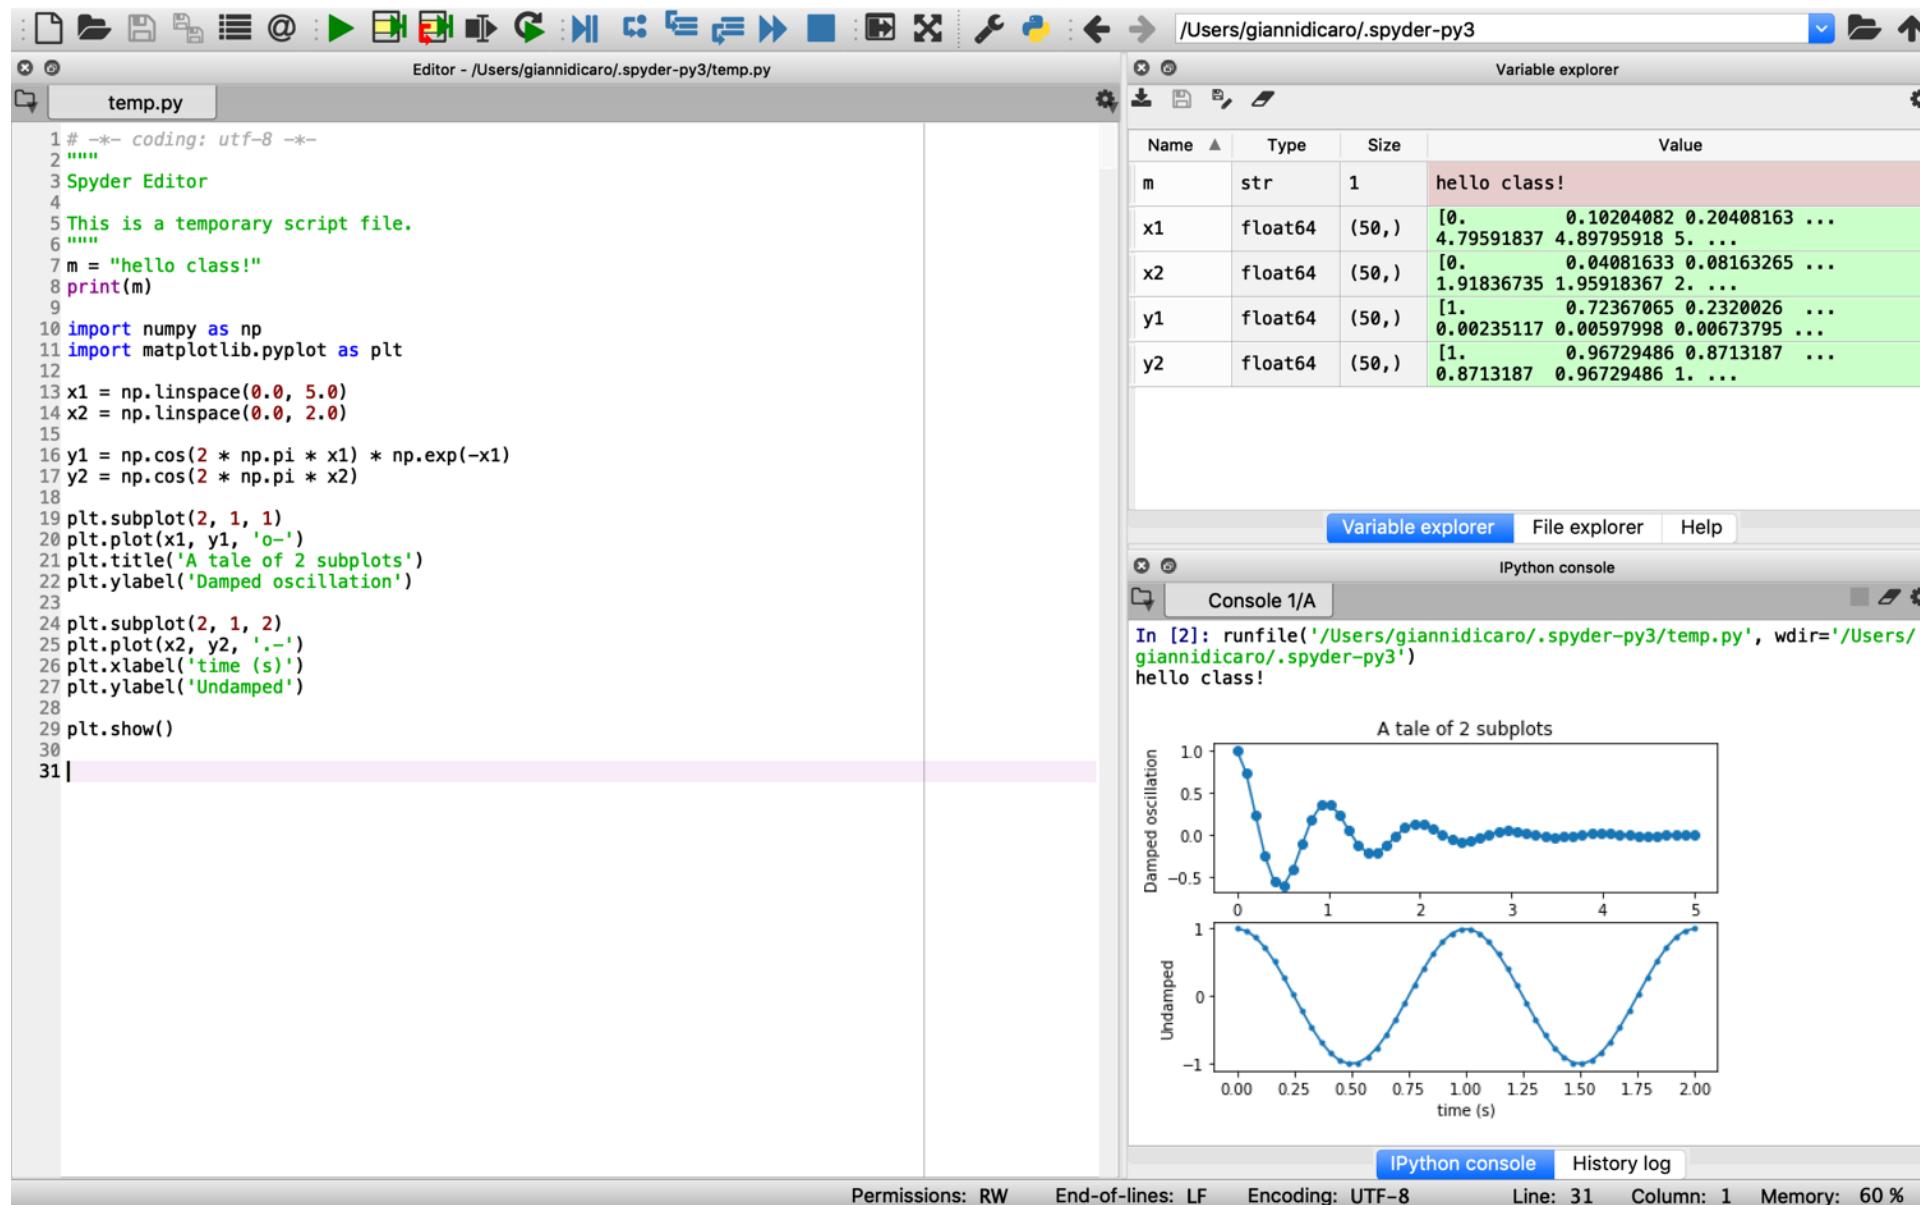


Python 3.8

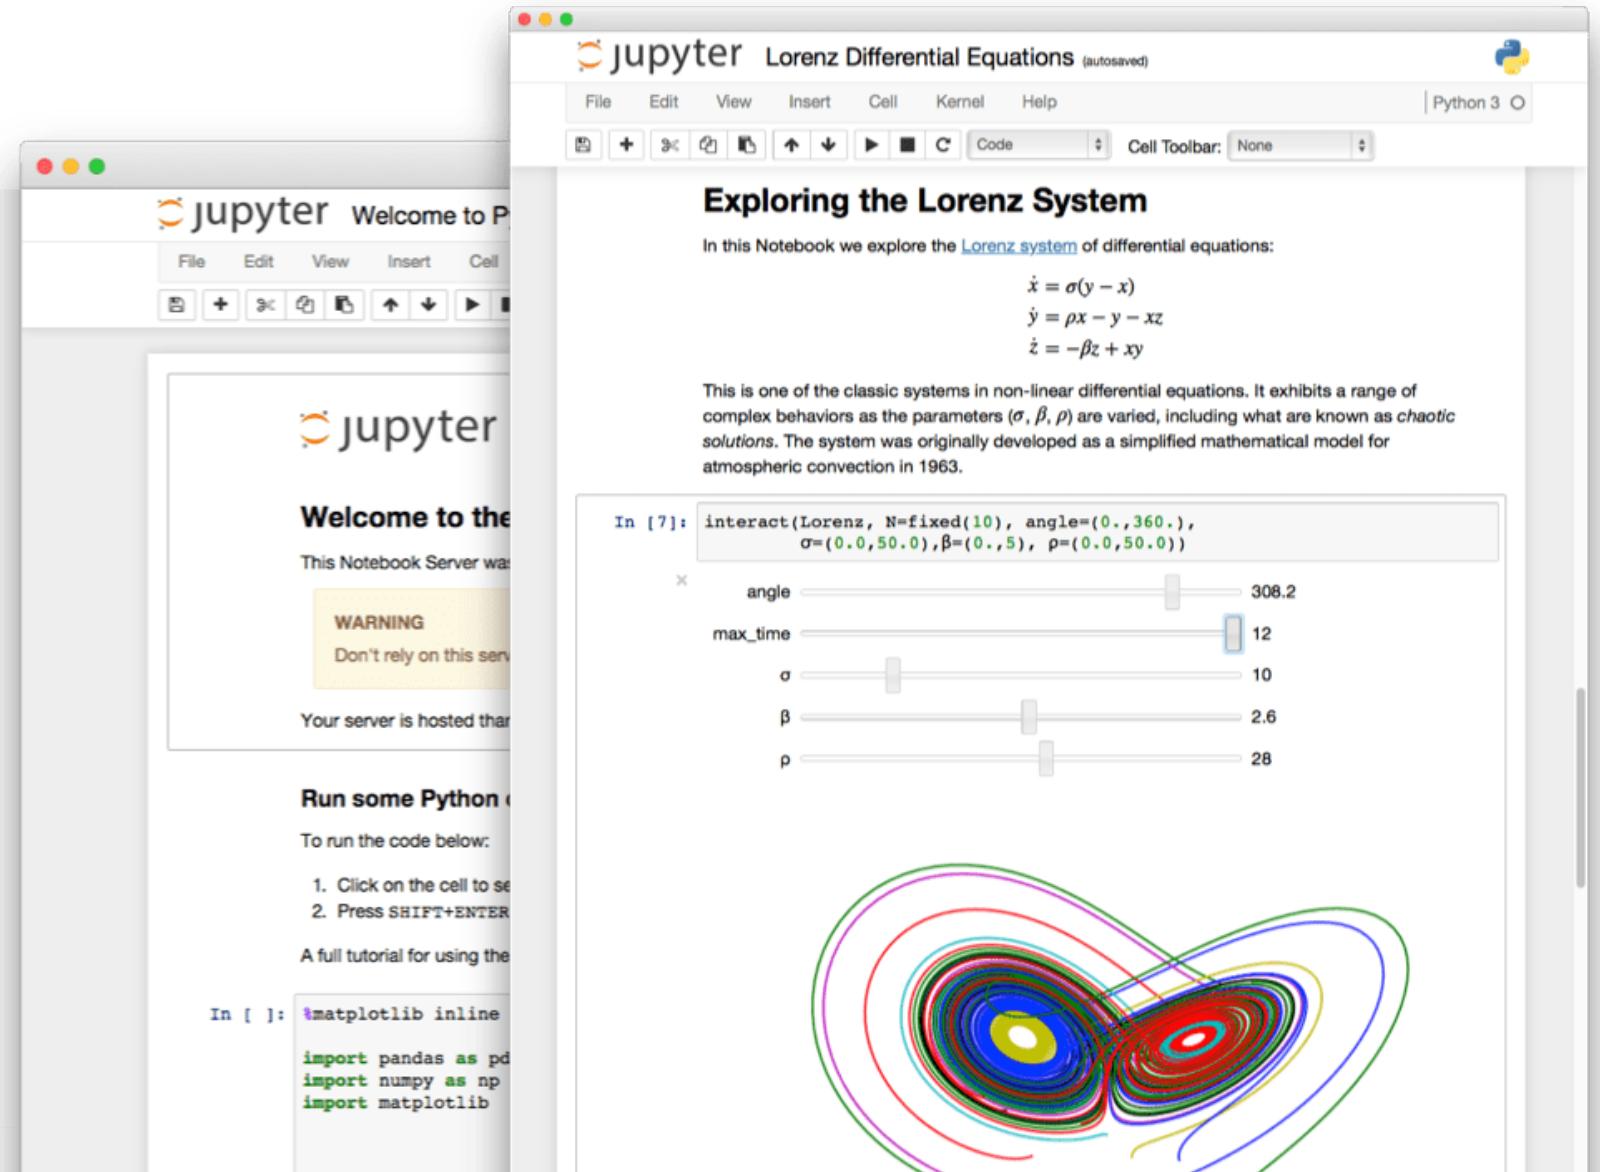
64-Bit (x86) Installer (529 MB)

64-Bit (Power8 and Power9) Installer (279 MB)

Spyder: Python Integrated Development Environment (IDE)



Jupyter Notebooks



Road map

- General overview: what this course is about
- Utility of the course
- Logistics: classes, labs, homework, exams, grading, website, books, piazza, rules
- Software for python programming
- More about the need to use efficient algorithms + computers: a practical example

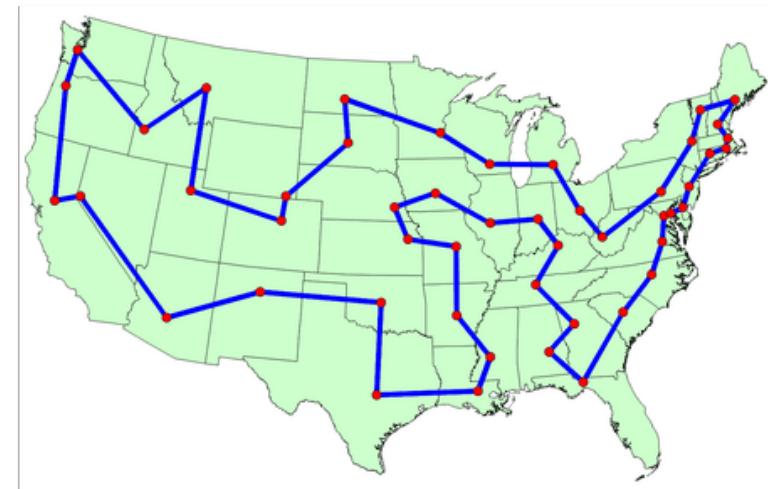
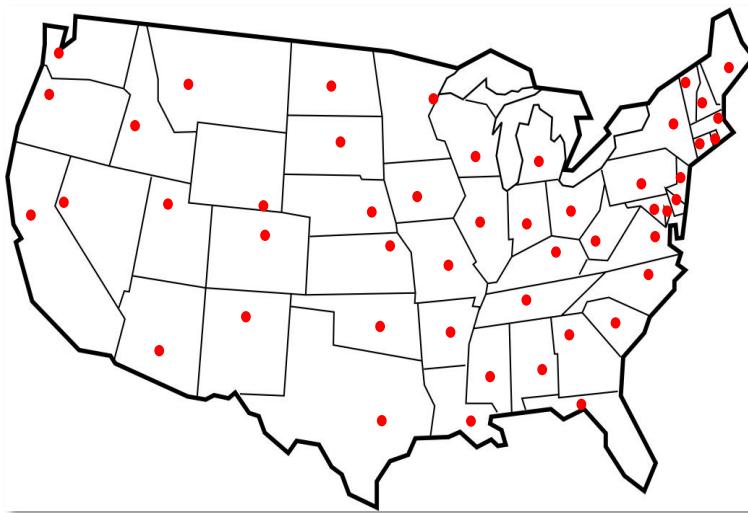
Are the (usual) problems really so difficult?



- Do we need efficient algorithms?
- Do we really need computers?
- Can we get the flavor of how difficult solving typical problems can be?
 - The following 10 slides illustrate these aspects considering one of the most important and well-studied problems in computer science, the **Traveling Salesman Problem** (TSP)

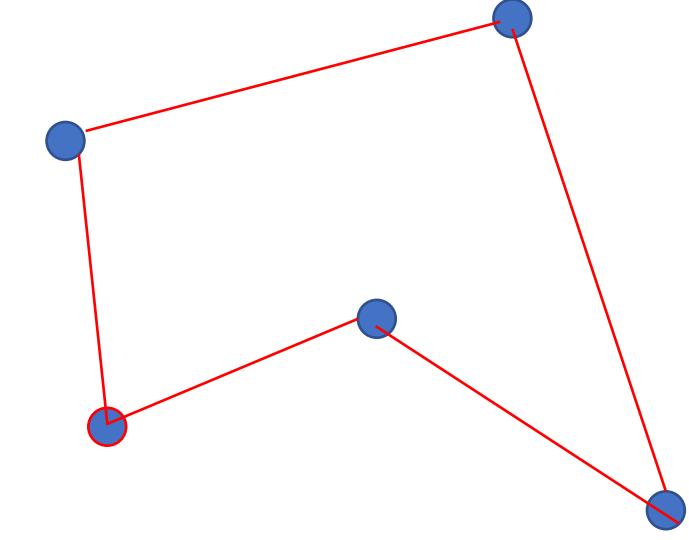
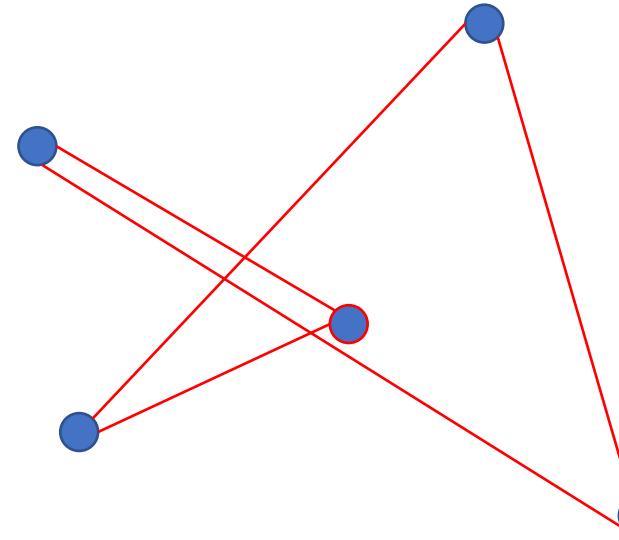
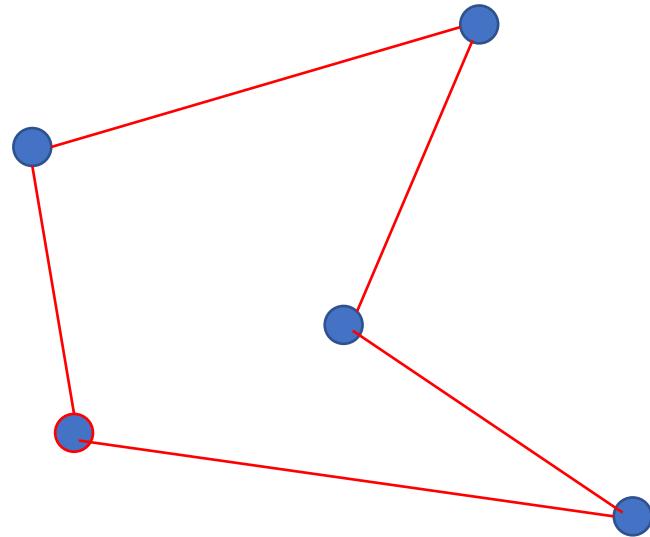
The (very famous) Traveling Salesman Problem (TSP)

- ✓ Given a set of n sites (e.g., cities) along with the cost of travel between each pair of them, the TSP is to find the cheapest way of visiting all the n sites once and only once, and returning to the starting point (**Closed tour**)
- ✓ Common cost measure is distance: Tour of shortest length

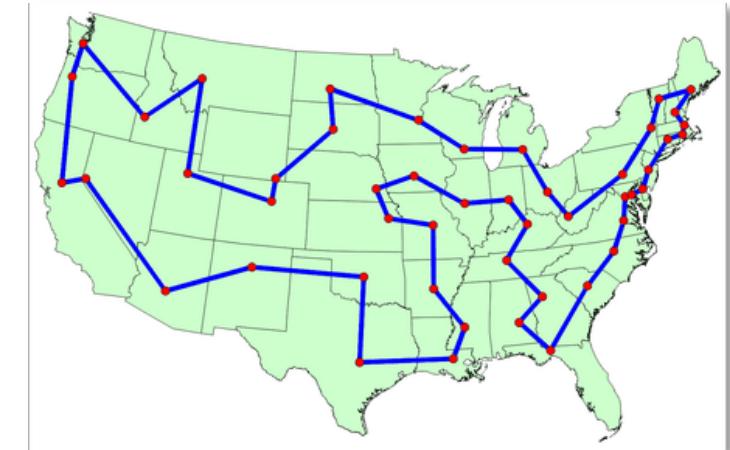
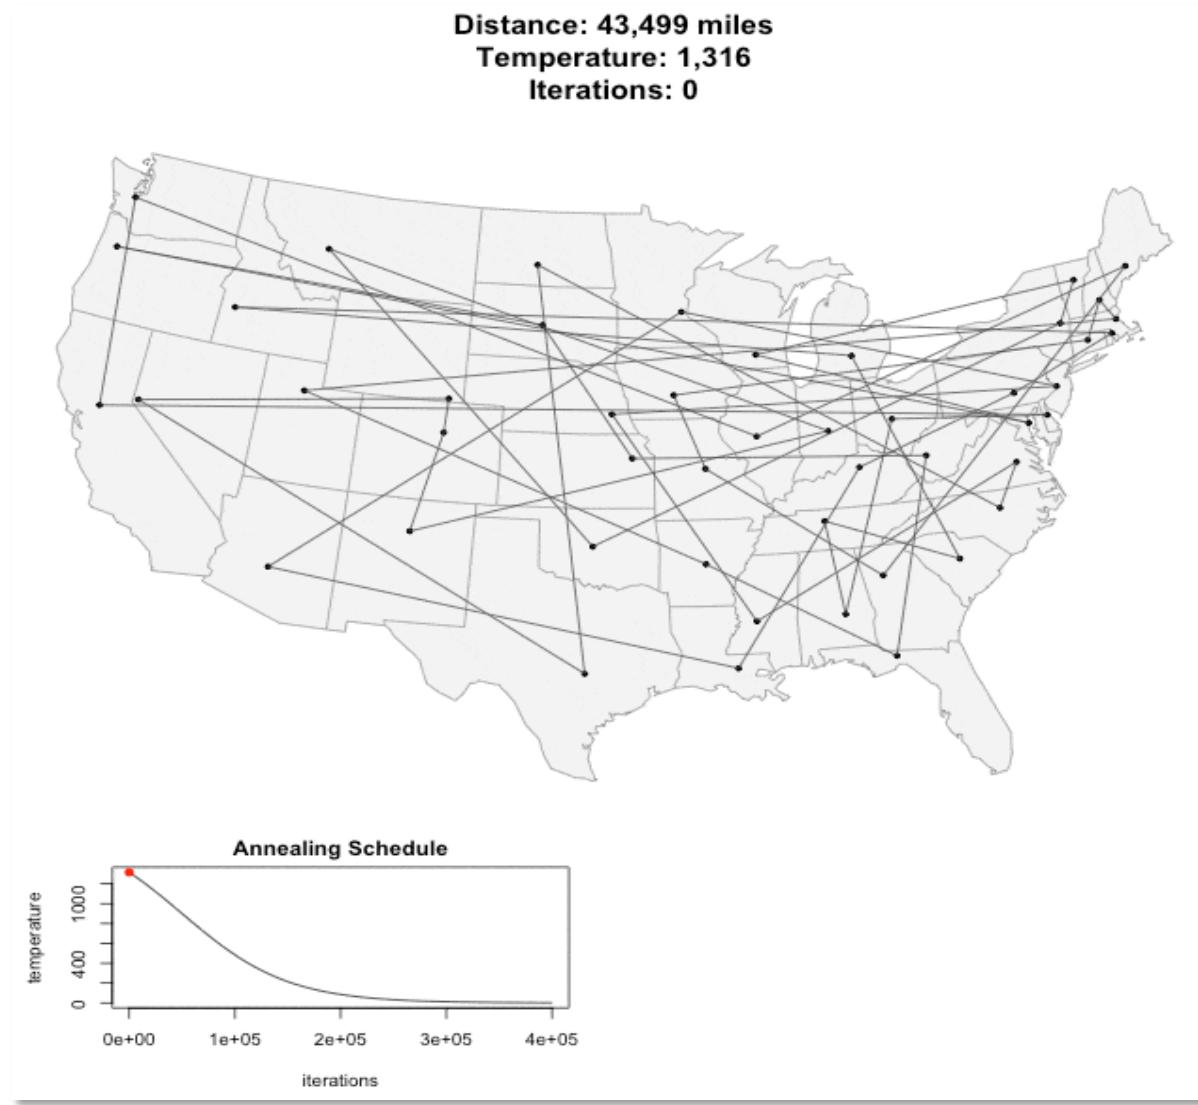


- ❖ The TSP is around since the 18th century, since it pertains to the practical optimization of the traveling paths of merchants, traders, postmen, ...

TSP: 5-sites, example of solutions with different costs



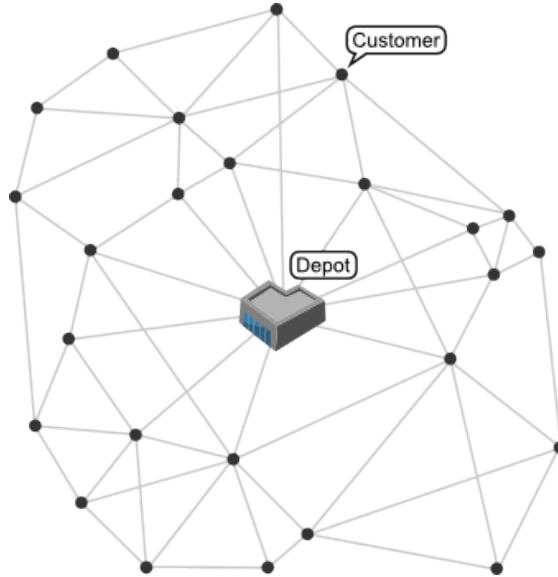
TSP among the capital cities of US states: many possible solutions!



The solution of minimal traveling distance!

TSP is an abstract model of many real-world problems!

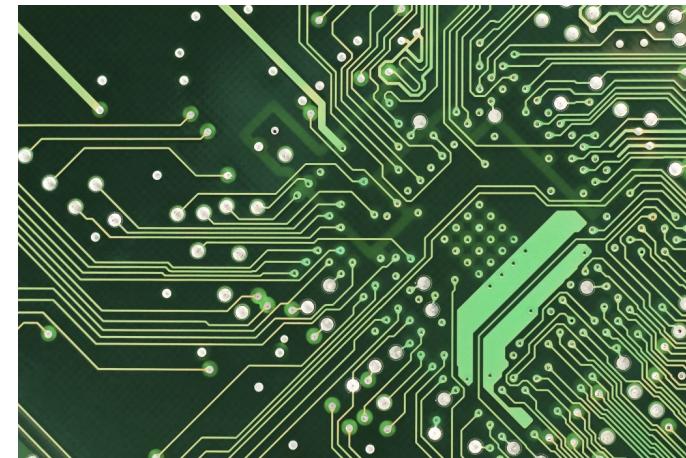
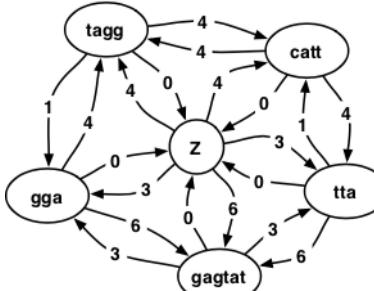
(Bus) route optimization



Delivery / Pick-up / Monitoring:

- Pizzas, Water
- Postal mail
- Express courier
- Flora, fauna
- ...

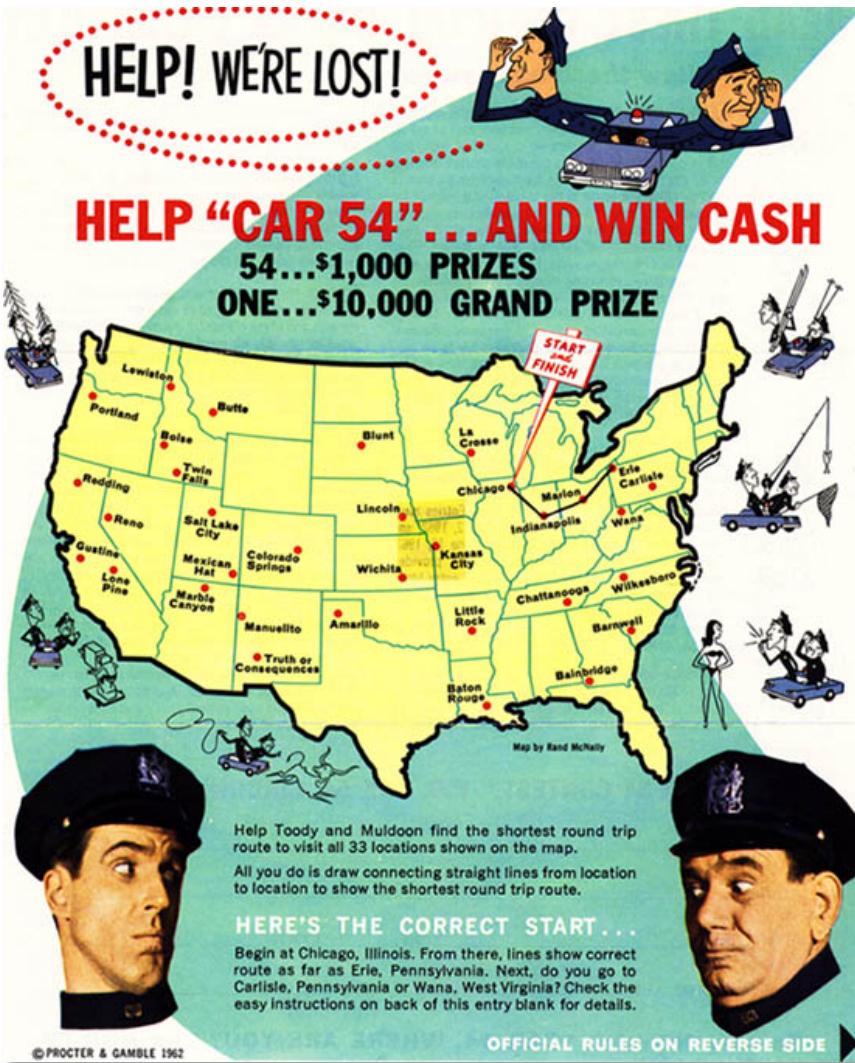
Genome sequencing



Microchips manufacturing

A big prize for solving a TSP instance, in 1962

- ✓ Procter & Gamble, 1962, 33-city contest in USA



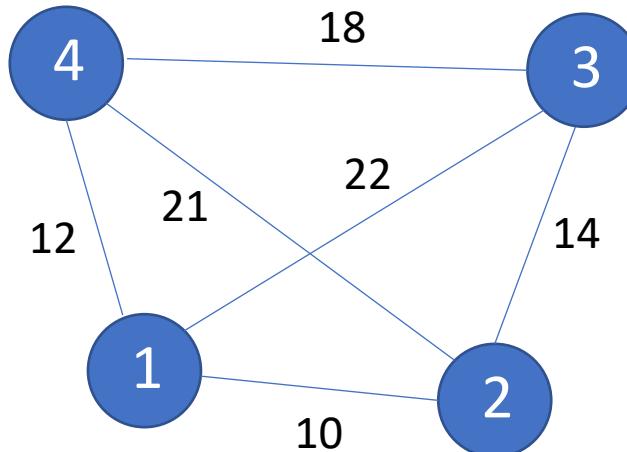
- 10,000 \$ (1962) → 500,000 \$ (2020)
 - Winner: Gerald Thompson, CMU! ☺

Is it so difficult?



How many feasible solutions are there? (we need the best one!)

- A TSP with $n = 4$ sites to visit: how many possible solutions?



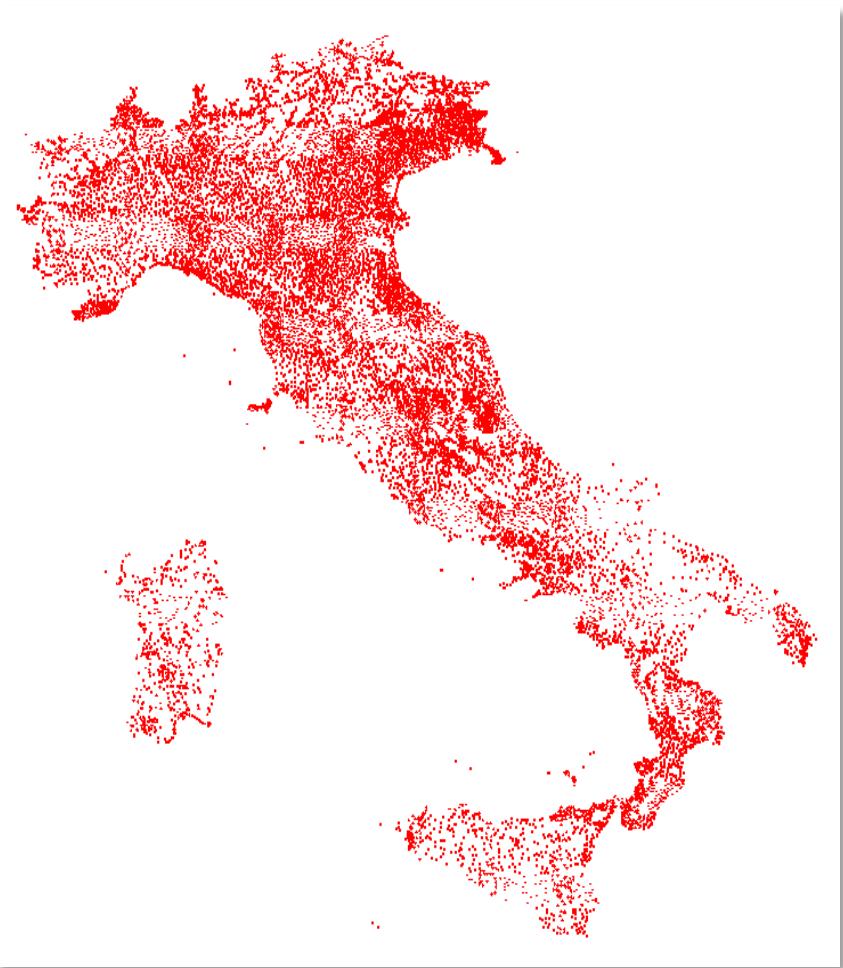
Optimal tour?

1234	2134	3124	4123
1243	2143	3142	4132
1324	2314	3214	4213
1342	2341	3241	4231
1423	2413	3412	4312
1432	2431	3421	4321

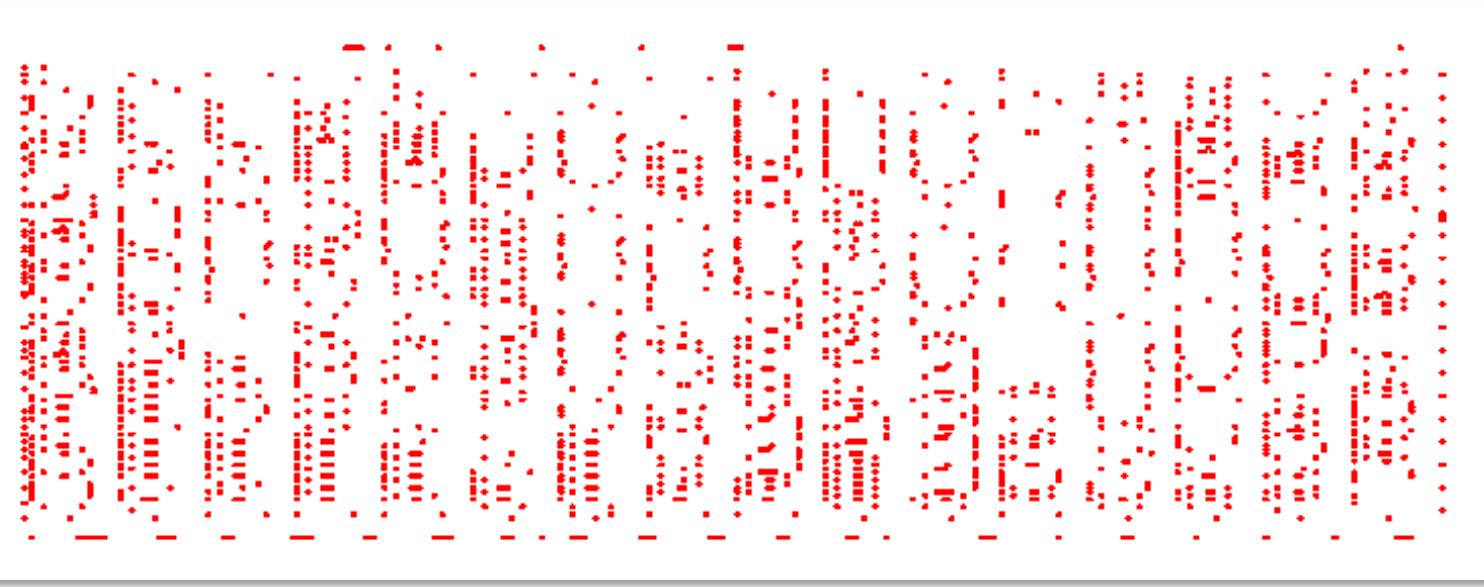
$$24 = 4 \cdot 3 \cdot 2 \cdot 1 = 4!$$

Number of feasible solutions: $\approx n!$

A very difficult problem when the number of sites is large



16,862 cities



21,144 sites

Optimal tour? By hand? → No way!

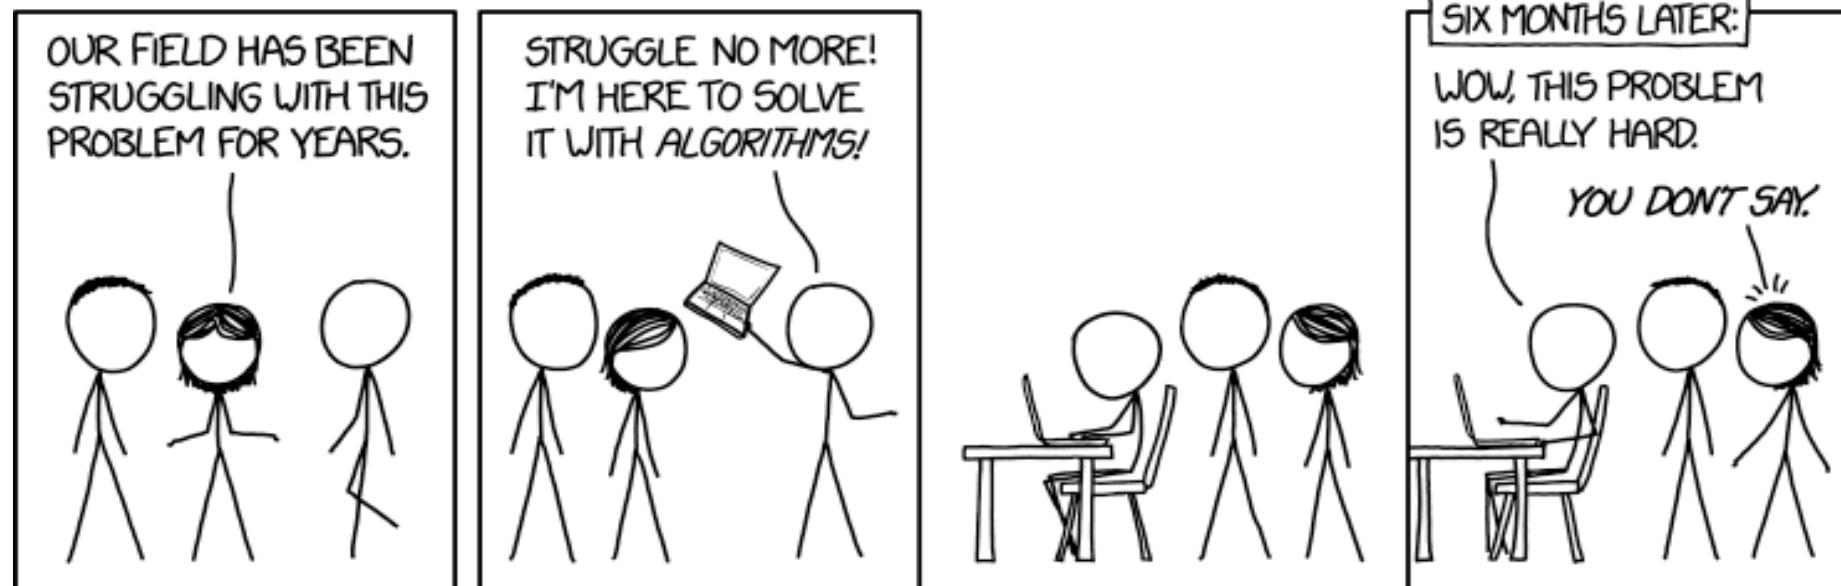
Factorial explosion, need for smart and efficient algorithms

Number of feasible solutions: $\approx n!$

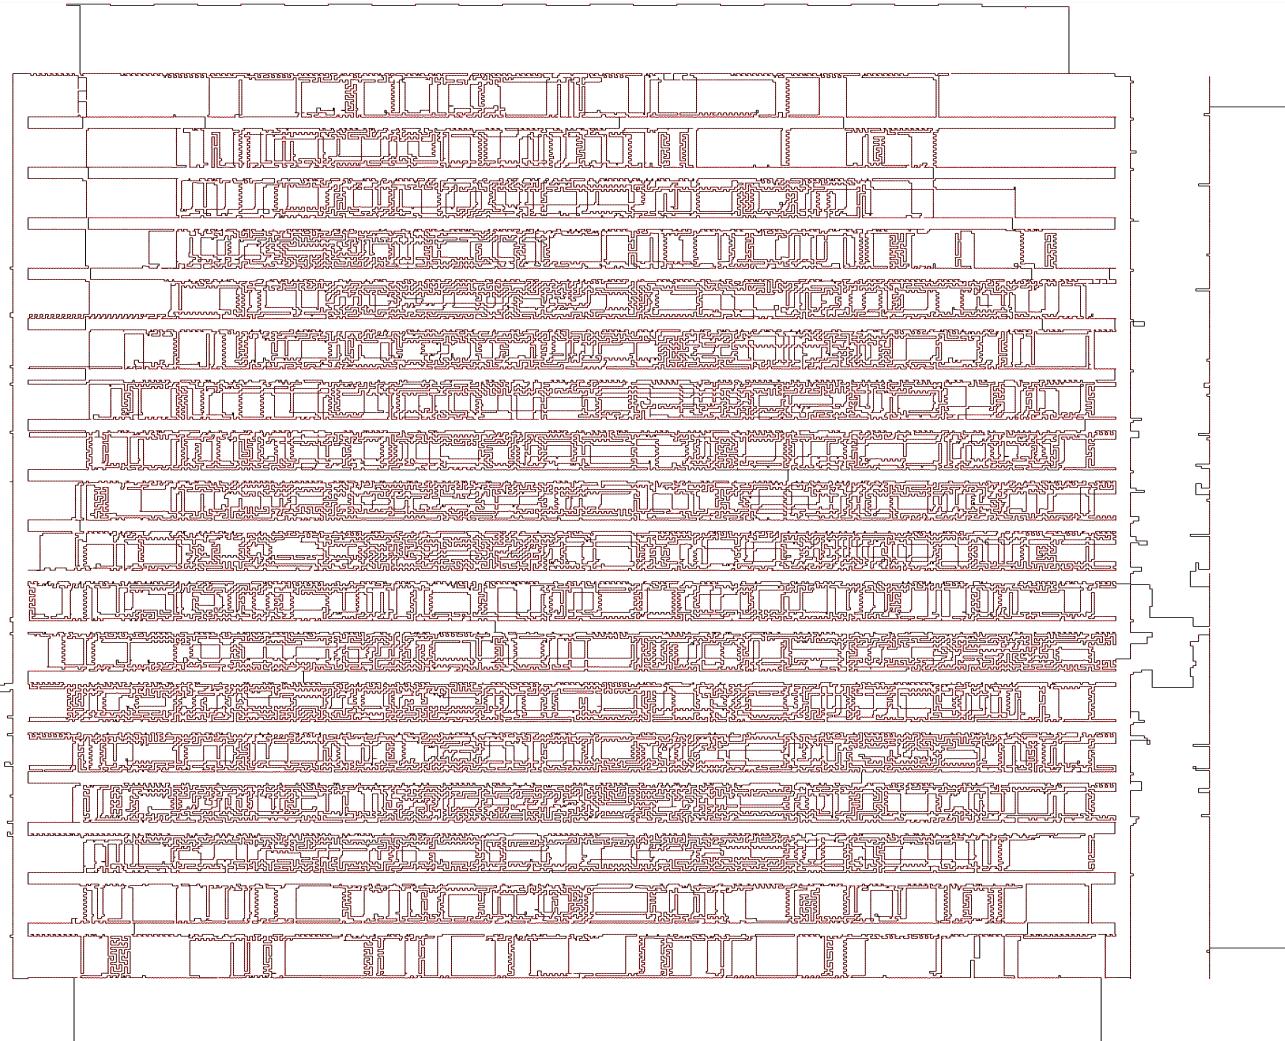
- $n = 100 \rightarrow 100! \approx 9 \cdot 10^{157}$
 - Number of atoms in the universe: $\approx 10^{80}$
 - Fastest computer in the world (IBM AC922): $\approx 10^{15}$ FP operations per sec
 - If it generates one solution per operation (not true): $\approx 10^{142}$ sec
 - $\approx 10^7$ sec in one year $\rightarrow \approx 10^{135}$ years!
 - Age of the universe: $\approx 10^9$ years ...
-
- ✓ Brute force (complete enumeration of all solutions) is not feasible!
 - ✓ We need to find out smart (and formally sound) ways to search the solution set
→ Intelligent & Efficient algorithms!!

NP-Hardness (just to let you know that such a thing does exist!)

- ✓ BTW, It doesn't matter how intelligent the algorithm is, in the worst case (i.e., for *some* problem instances) the time to find the optimal solution for a given TSP instance will **explode exponentially** with n 😞
- ✓ NP-hardness (conjecture)



Largest solved (with guarantees of optimality of the solution found)



- Largest TSP instance solved to optimality with guarantees: 85,900 sites, from VLSI

The algorithm is smart: it exploits the structure of the problem for searching among all the possible solutions without searching among all of them 😊

Summary

- **Overview of the topics:** a journey into computation
 - Problem solving
 - Designing and understanding algorithms
 - Fundamentals programming constructs
 - Core techniques for computational problem-solving
 - Python as programming language
 - Tools (graphics, data manipulation, data publishing)
 - Applications
- **Action strategies:** first think, reason, abstract and model, design the algorithm, implement it by starting with a little chunk of code for a part of it, test it out, revise it, keep adding pieces of code, test with different inputs, understand what you've done!
- **The good student:** Actively participates to lectures + READs the slides + READs technical notes + READs the book chapters + Asks questions + Tries out code, tries out code, tries out code
.... ☺